

```
In [7]: %pip install numpy
```

```
Requirement already satisfied: numpy in /opt/anaconda3/envs/datateq/lib/python3.13/site-packages (2.3.1)
Note: you may need to restart the kernel to use updated packages.
```

```
In [8]: iteration = 0
W_0 = 0
W_1 = 0
eta = 0.001 # step size
conv_thresh = 0.01

X = [1, 2, 4, 6, 8]
Y = [2, 5, 6, 9, 11]

weights = [W_0, W_1]
```

```
In [9]: import math

# def get_gradients():
#     pass

def gradient_desc(a, b, eta, X, Y, conv_thresh, iteration):
    converged = False
    while not converged:

        #get partial derivative of a and b
        grad_a = -2 * sum((Y[i] - (a + b*X[i])) for i in range(len(X)))
        grad_b = -2 * sum((Y[i] - (a + b * X[i])) * X[i] for i in range(len(X)))

        #get magnitude of gradient
        grad_mag = math.sqrt((grad_a)**2 + (grad_b)**2)

        # check if we have convergence
        if grad_mag < conv_thresh:
            converged = True
            print("=*60")
            print(f"Converged!")
            print(f"iteration: {iteration} \n    gradient magnitude: {grad_mag}")

        if iteration % 100 == 0:
            print("=*60")
            print(f"iteration: {iteration} \n    gradient magnitude: {grad_mag}")

        # update coefficients
        a = a - eta * grad_a
        b = b - eta * grad_b

        # iterate count
        iteration += 1

    return a, b, iteration
```

```
In [10]: final_a, final_b, iterations = gradient_desc(W_0, W_1, eta, X, Y, conv_thres

print("=*60)
print(f"""
final a: {final_a}
final b: {final_b}
iterations: {iterations}
""")
```

```
=====
iteration: 0
    gradient magnitude: 362.0662922725616
=====
iteration: 100
    gradient magnitude: 2.676170860560416
=====
iteration: 200
    gradient magnitude: 2.0564409199245373
=====
iteration: 300
    gradient magnitude: 1.5802239384126953
=====
iteration: 400
    gradient magnitude: 1.2142861345241884
=====
iteration: 500
    gradient magnitude: 0.9330897859823554
=====
iteration: 600
    gradient magnitude: 0.7170110272615112
=====
iteration: 700
    gradient magnitude: 0.5509703577704025
=====
iteration: 800
    gradient magnitude: 0.4233802878891129
=====
iteration: 900
    gradient magnitude: 0.3253366821736797
=====
iteration: 1000
    gradient magnitude: 0.24999736595081742
=====
iteration: 1100
    gradient magnitude: 0.19210463008589795
=====
iteration: 1200
    gradient magnitude: 0.14761831093731925
=====
iteration: 1300
    gradient magnitude: 0.1134338392273204
=====
iteration: 1400
    gradient magnitude: 0.08716558128966297
=====
iteration: 1500
    gradient magnitude: 0.06698035271766711
=====
iteration: 1600
    gradient magnitude: 0.05146948582003237
=====
iteration: 1700
    gradient magnitude: 0.03955052284876008
=====
iteration: 1800
```

```

gradient magnitude: 0.03039167445891953
=====
iteration: 1900
    gradient magnitude: 0.02335377157841935
=====
iteration: 2000
    gradient magnitude: 0.01794565967973375
=====
iteration: 2100
    gradient magnitude: 0.013789922551025787
=====
iteration: 2200
    gradient magnitude: 0.010596543529579961
=====
Converged!
iteration: 2222
    gradient magnitude: 0.009999925267377408
=====

final a: 1.5511437620712774
final b: 1.2018747340626836
iterations: 2223

```

```
In [11]: def get_ground_truth(X, Y):
    sum_x = sum(X)
    sum_y = sum(Y)
    sum_xy = sum(X[i] * Y[i] for i in range(len(X)))
    sum_x2 = sum(X[i]**2 for i in range(len(X)))
    n = len(X)

    #solve grad for θ
    b_truth = (n * sum_xy - sum_x * sum_y) / (n * sum_x2 - sum_x**2)
    a_truth = a = (sum_y - b_truth * sum_x) / n

    return a_truth, b_truth

def eval(model, X, Y):
    a_truth, b_truth = get_ground_truth(X, Y)
    a, b = model
    RSS_gd = sum((Y[i] - (a + b*X[i]))**2 for i in range(len(X)))
    RSS_truth = sum((Y[i] - (a_truth + b_truth*X[i]))**2 for i in range(len(X)))

    diff_RSS = abs(RSS_gd - RSS_truth)

    print(f"RSS from Gradient Descent: {RSS_gd}")
    print(f"RSS from Analytical: {RSS_truth}")
    print(f"Difference: {diff_RSS}")


In [12]: eval((final_a, final_b), X, Y)
```

```
RSS from Gradient Descent: 1.8719701261724233
RSS from Analytical: 1.871951219512195
Difference: 1.8906660228301675e-05
```

1. How many iterations do you have for convergence? **2223**
2. What are parameter values in your best fitted simple linear regression model?
final a: 1.5511437620712774
final b: 1.2018747340626836